



References

A Governance Layer for AI Development

Technical Product Dossier

Luis López Trejo

Sep 2026

References

Sources and Supporting Material

Amazon Web Services. (n.d.). *Verify feature consistency across training and inference*. AWS Well-Architected Machine Learning Lens. Retrieved September 10, 2026, from <https://docs.aws.amazon.com/wellarchitected/latest/machine-learning-lens/mlrel03-bp02.html>

AWS-FEAT

Association for the Advancement of Artificial Intelligence. (n.d.). *AAAI reproducibility checklist*. Retrieved September 10, 2026, from <https://aaai.org/conference/aaai/aaai-26/reproducibility-checklist/>

AAAI-CHK

Breck, E., Cai, S., Nielsen, E., Salib, M., & Sculley, D. (2017). The ML test score: A rubric for ML production readiness and technical debt reduction. In *2017 IEEE International Conference on Big Data (Big Data)* (pp. 1123–1132). IEEE. <https://doi.org/10.1109/BigData.2017.8258038>

BRE17

Bóveda 2.0—S5 Signals foundation. (n.d.). (Version 0.1) [Unpublished internal design document].

BOV-S5-FOUNDATION

Bóveda 2.0—Supervisory Backbone. (n.d.). (Version 0.2) [Unpublished internal design document].

BOV-BACKBONE

Bóveda S5 Alpha Signal Catalogue. (n.d.). (Version 1.0-alpha) [Unpublished internal implementation document].

BOV-S5-CATALOGUE

Gebru, T., Morgenstern, J., Vecchione, B., Vaughan, J. W., Wallach, H., Daumé III, H., & Crawford, K. (2021). Datasheets for datasets. *Communications of the ACM*, 64(12), 86–92. <https://doi.org/10.1145/3458723>

GEB21

Kapoor, S., & Narayanan, A. (2023). Leakage and the reproducibility crisis in machine-learning-based science. *Patterns*, 4(9), Article 100804. <https://doi.org/10.1016/j.patter.2023.100804>

KAP23

Kaufman, S., Rosset, S., Perlich, C., & Stitelman, O. (2012). Leakage in data mining: Formulation, detection, and avoidance. *ACM Transactions on Knowledge Discovery from Data*, 6(4), Article 15, 1–21. <https://doi.org/10.1145/2382577.2382579>

KAU12

Mitchell, M., Wu, S., Zaldivar, A., Barnes, P., Vasserman, L., Hutchinson, B., Spitzer, E., Raji, I. D., & Gebru, T. (2019). Model cards for model reporting. In *Proceedings of the Conference on Fairness, Accountability, and Transparency* (pp. 220–229). Association for Computing Machinery. <https://doi.org/10.1145/3287560.3287596>

MIT19

Neural Information Processing Systems Foundation. (n.d.). *Paper checklist*. Retrieved September 10, 2026, from <https://neurips.cc/public/guides/PaperChecklist>

NEU-CHK

Organisation for Economic Co-operation and Development. (2024). *Recommendation of the Council on Artificial Intelligence* (OECD/LEGAL/0449). <https://legalinstruments.oecd.org/en/instruments/OECD-LEGAL-0449>

OECD24

References

Sources and Supporting Material

Papers with Code. (n.d.). *Machine learning code completeness checklist*. GitHub. Retrieved September 10, 2026, from <https://github.com/paperswithcode/releasing-research-code>

CODE-COMP

Pineau, J., Vincent-Lamarre, P., Sinha, K., Larivière, V., Beygelzimer, A., d'Alché-Buc, F., Fox, E., & Larochelle, H. (2021). Improving reproducibility in machine learning research: A report from the NeurIPS 2019 reproducibility program. *Journal of Machine Learning Research*, 22(164), 1–20. <https://www.jmlr.org/papers/v22/20-303.html>

PIN21

scikit-learn developers. (n.d.-a). *Common pitfalls and recommended practices*. Retrieved September 10, 2026, from https://scikit-learn.org/stable/common_pitfalls.html

SKL-PIT

scikit-learn developers. (n.d.-b). *Probability calibration*. Retrieved September 10, 2026, from <https://scikit-learn.org/stable/modules/calibration.html>

SKL-CAL

scikit-learn developers. (n.d.-c). *TimeSeriesSplit*. Retrieved September 10, 2026, from https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.TimeSeriesSplit.html

SKL-TS

Sculley, D., Holt, G., Golovin, D., Davydov, E., Phillips, T., Ebner, D., Chaudhary, V., Young, M., Crespo, J.-F., & Dennison, D. (2015). Hidden technical debt in machine learning systems. In C. Cortes, N. Lawrence, D. Lee, M. Sugiyama, & R. Garnett (Eds.), *Advances in neural information processing systems* (Vol. 28, pp. 2503–2511). Curran Associates, Inc. <https://proceedings.neurips.cc/paper/2015/hash/86df7dcfd896fcdf2674f757a2463eba-Abstract.html>

SCU15

Tabassi, E. (2023). *Artificial intelligence risk management framework (AI RMF 1.0)* (NIST AI 100-1). National Institute of Standards and Technology. <https://doi.org/10.6028/NIST.AI.100-1>

NIST23

Varma, S., & Simon, R. (2006). Bias in error estimation when using cross-validation for model selection. *BMC Bioinformatics*, 7, Article 91. <https://doi.org/10.1186/1471-2105-7-91>

VAR06

Zinkevich, M. (n.d.). *Rules of machine learning: Best practices for ML engineering*. Google for Developers. Retrieved September 10, 2026, from <https://developers.google.com/machine-learning/guides/rules-of-ml/>

ZIN-RML